


▼ Tutorial: Object Detection with TFLite


Introduction

Imagine that you have a niece or a nephew and you want to give them a present. When you were growing up, your aunt gave you a "Find the Duck" book. You had lots of fun finding the duck on every page of this board book. Today, you want to make this book into a computer game. For that, you need to be able to teach the computer how to find the duck. This is what this tutorial will teach you.



Our plan

The task that you are about to undertake is called "Object Detection." The good news is that the Google library called TensorFlow already does most of the groundwork for object detection. Furthermore, the TensorFlow Lite part of the library will help you to put your application on a phone or a device app. The end result of your object detection will look like a screenshot below, where you will be able to detect, out of a known set of objects, which ones are present in our picture and what are their locations.



We will do it in three steps. First, you will have to prepare the data: those objects that you will be looking to identify. After you got the objects, you will have to convert them to TFRecord format that Object Detection API expects. Then, you will train the model with this data. And finally, you will export the model to TFLite, preparing it to be used in your phone app. In the next tutorial, we will teach you how to use the resulting TFLite model in your phone app. So, let us start.

▼ Data collection

We have taken a dataset with pictures of fruit. Each data sample has an images of a fruit and an XML file with the fruit coordinates. This dataset came from Kaggle [here](#), with a Creative Commons license, so for ease of use we have placed it in next to this tutorial.

In Jupyter Notebook, you can do command lines if you only start with a exclamation sign. Let us download this dataset first.

```
!wget -nc https://github.com/elephantscale/E2E-Object-Detection-in-TFLite/raw/master/  
File 'Fruit_Images_for_Object_Detection.zip' already there; not retrieving.
```

In the same way, let us unzip the dataset

```
!unzip -qqn Fruit_Images_for_Object_Detection.zip
```

▼ Generate intermediate files

To be able to generate TFRecords from our fruits dataset we first generate a .csv file that would contain the following fields -

- filename
- width
- height
- class
- xmin
- ymin
- xmax
- ymax

Now, we need to convert the XML descriptions to CSV. In your case, you may have to adjust this code, depending on the format of the data in your real-life dataset.

```
# Convert XML to CSV
```

```
import os
import glob
import pandas as pd
import xml.etree.ElementTree as ET

def xml_to_csv(path):
    xml_list = []
    for xml_file in glob.glob(path + '/*.xml'):
        tree = ET.parse(xml_file)
        root = tree.getroot()
        for member in root.findall('object'):
            value = (root.find('filename').text,
                    int(root.find('size')[0].text),
                    int(root.find('size')[1].text),
                    member[0].text,
                    int(member[4][0].text),
                    int(member[4][1].text),
                    int(member[4][2].text),
                    int(member[4][3].text)
                    )
            xml_list.append(value)
    column_name = ['filename', 'width', 'height', 'class', 'xmin', 'ymin', 'xmax', 'ymin', 'ymax']
    xml_df = pd.DataFrame(xml_list, columns=column_name)
```

```
return xml_df
```

```
def call_xml_to_csv():  
    train = "/content/train_zip/train"  
    test = "/content/test_zip/test"  
    for directory in [train, test]:  
        xml_df = xml_to_csv(directory)  
        xml_df.to_csv('{}_labels.csv'.format(directory), index=None)  
        print('Successfully converted xml to csv.')
```

```
call_xml_to_csv()
```

```
Successfully converted xml to csv.  
Successfully converted xml to csv.
```

```
!head -5 /content/train_zip/train_labels.csv
```

```
filename,width,height,class,xmin,ymin,xmax,ymax  
banana_8.jpg,1920,1280,banana,496,168,1603,923  
banana_8.jpg,1920,1280,banana,199,587,1470,1268  
banana_8.jpg,1920,1280,banana,797,1,1772,611  
banana_8.jpg,1920,1280,banana,1052,1,1916,550
```

```
!head -5 /content/test_zip/test_labels.csv
```

```
filename,width,height,class,xmin,ymin,xmax,ymax  
banana_83.jpg,630,355,banana,63,17,560,323  
orange_81.jpg,1300,990,orange,90,287,608,809  
orange_81.jpg,1300,990,orange,560,217,1145,769  
orange_81.jpg,1300,990,orange,323,37,786,551
```

Now that we have .csv files we can do some basic exploratory data analysis (EDA) to better understand the dataset.

▼ Basic EDA

```
import matplotlib.pyplot as plt  
import pandas as pd  
import numpy as np  
import cv2  
import os
```

```
train_df = pd.read_csv("/content/train_zip/train_labels.csv")  
test_df = pd.read_csv("/content/test_zip/test_labels.csv")
```

```
train_df.head()
```

	filename	width	height	class	xmin	ymin	xmax	ymax
0	banana_8.jpg	1920	1280	banana	496	168	1603	923
1	banana_8.jpg	1920	1280	banana	199	587	1470	1268
2	banana_8.jpg	1920	1280	banana	797	1	1772	611
3	banana_8.jpg	1920	1280	banana	1052	1	1916	550
4	apple_31.jpg	780	439	apple	304	105	773	439

```
test_df.head()
```

	filename	width	height	class	xmin	ymin	xmax	ymax
0	banana_83.jpg	630	355	banana	63	17	560	323
1	orange_81.jpg	1300	990	orange	90	287	608	809
2	orange_81.jpg	1300	990	orange	560	217	1145	769
3	orange_81.jpg	1300	990	orange	323	37	786	551
4	orange_77.jpg	732	549	orange	1	123	130	299

```
train_df["class"].value_counts()
```

```
banana    169
apple     156
orange    140
Name: class, dtype: int64
```


```
test_df["class"].value_counts()
```

```
orange     42
banana     40
apple      35
Name: class, dtype: int64
```


```
def show_images(df, is_train=True):
    if is_train:
        root = "/content/train_zip/train"
    else:
        root = "/content/test_zip/test"
    plt.figure(figsize=(15,15))
    for i in range(10):
        n = np.random.choice(df.shape[0], 1)
        plt.subplot(5,5,i+1)
        plt.xticks([])
```

```
plt.yticks([])
plt.grid(True)
image = plt.imread(os.path.join(root, df["filename"][int(n)]))
plt.imshow(image)
label = df["class"][int(n)]
plt.xlabel(label)
plt.show()
```

```
show_images(train_df)
```



```
show_images(test_df, is_train=False)
```



```
def verify_annotations(df, is_train=True):
```

```

if is_train:
    root = "/content/train_zip/train"
else:
    root = "/content/test_zip/test"



plt.figure(figsize=(12,12))
for i in range(3):
    n = np.random.choice(df.shape[0], 1)
    plt.subplot(1,3,i+1)
    plt.xticks([])
    plt.yticks([])

    image = plt.imread(os.path.join(root, df["filename"][int(n)]))
    xmin, ymin = int(df["xmin"][int(n)]), int(df["ymin"][int(n)])
    xmax, ymax = int(df["xmax"][int(n)]), int(df["ymax"][int(n)])
    cv2.rectangle(image, (xmin, ymin), (xmax, ymax), (255,0,0), 3)
    plt.imshow(image)




plt.show()

```

```
verify_annotations(train_df, is_train=True)
```



```
verify_annotations(test_df, is_train=False)
```



As we can see the dataset has annotation issues. So, our model training can suffer a lot from this.

▼ Generate TFRecords and .pbtxt

The TFRecord format is a simple format for storing a sequence of binary records. The format is explained [here](#) but for our purposes it is enough that the code below created these records for us.

The utility scripts that I used in the following cells were adapted from [this repository](#).

TODO - should we convert to TF2 here?

```
%tensorflow_version 1.x
import tensorflow as tf
print(tf.__version__)
```

```
!git clone https://github.com/tensorflow/models.git
```

```
% cd models/research
!pip install --upgrade pip
# Compile protos.
!protoc object_detection/protos/*.proto --python_out=.
# Install TensorFlow Object Detection API.
!cp object_detection/packages/tf1/setup.py .
!python -m pip install --use-feature=2020-resolver .
```

1.15.2

Cloning into 'models'...

remote: Enumerating objects: 49553, done.

remote: Total 49553 (delta 0), reused 0 (delta 0), pack-reused 49553

Receiving objects: 100% (49553/49553), 558.66 MiB | 39.40 MiB/s, done.

Resolving deltas: 100% (34186/34186), done.

/content/models/research/models/research

Requirement already satisfied: pip in /usr/local/lib/python3.6/dist-packages (20

WARNING: --use-feature=2020-resolver no longer has any effect, since it is now t

Processing /content/models/research/models/research

Requirement already satisfied: pillow in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: lxml in /usr/local/lib/python3.6/dist-packages (f

Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packa

Requirement already satisfied: Cython in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: contextlib2 in /usr/local/lib/python3.6/dist-pack

Requirement already satisfied: tf-slim in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (fr

Requirement already satisfied: pycocotools in /usr/local/lib/python3.6/dist-pack

Requirement already satisfied: lvis in /usr/local/lib/python3.6/dist-packages (f

Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (

Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: python-dateutil>=2.8.0 in /usr/local/lib/python3.

Requirement already satisfied: kiwisolver>=1.1.0 in /usr/local/lib/python3.6/dis

Requirement already satisfied: pyparsing>=2.4.0 in /usr/local/lib/python3.6/dist


Requirement already satisfied: numpy>=1.18.2 in /usr/local/lib/python3.6/dist-pa

Requirement already satisfied: cycloper>=0.10.0 in /usr/local/lib/python3.6/dist-p

Requirement already satisfied: opencv-python>=4.1.0.25 in /usr/local/lib/python3

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-pac


```
Requirement already satisfied: setuptools>=18.0 in /usr/local/lib/python3.6/dist
Requirement already satisfied: absl-py>=0.2.2 in /usr/local/lib/python3.6/dist-p
Building wheels for collected packages: object-detection
  Building wheel for object-detection (setup.py) ... done
  Created wheel for object-detection: filename=object_detection-0.1-py3-none-any
  Stored in directory: /tmp/pip-ephem-wheel-cache-12tt468k/wheels/ab/a8/ef/cead6
Successfully built object-detection
Installing collected packages: object-detection
  Attempting uninstall: object-detection
    Found existing installation: object-detection 0.1
    Uninstalling object-detection-0.1:
      Successfully uninstalled object-detection-0.1
Successfully installed object-detection-0.1
```



```
#!/wget https://raw.githubusercontent.com/elephantscale/E2E-Object-Detection-in-TFLite
!wget https://raw.githubusercontent.com/elephantscale/E2E-Object-Detection-in-TFLite/
```

```
--2020-12-25 20:26:57-- https://raw.githubusercontent.com/elephantscale/E2E-Obj
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.0.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.0.13
HTTP request sent, awaiting response... 200 OK
Length: 3740 (3.7K) [text/plain]
Saving to: 'generate_tfrecord.py'
```


```
generate_tfrecord.p 100%[=====>] 3.65K --.-KB/s in 0s
```

```
2020-12-25 20:26:57 (74.7 MB/s) - 'generate_tfrecord.py' saved [3740/3740]
```



```
!python generate_tfrecord.py \
--csv_input=/content/train_zip/train_labels.csv \
--output_path=/content/train_zip/train.record

WARNING:tensorflow:From generate_tfrecord.py:107: The name tf.app.run is depreca
WARNING:tensorflow:From generate_tfrecord.py:88: The name tf.python_io.TFRecordW
W1225 20:26:59.525310 140510281471872 module_wrapper.py:139] From generate_tfrec
WARNING:tensorflow:From generate_tfrecord.py:47: The name tf.gfile.GFile is depr
W1225 20:26:59.573870 140510281471872 module_wrapper.py:139] From generate_tfrec
Successfully created the TFRecords: /content/train_zip/train.record
```



Before the running the cell below please edit the path variable in the main() function of generate_tfrecord.py. generate_tfrecord.py should be located here - </content/models/research>.

```
!python generate_tfrecord.py \
```

```
!python generate_tfrecord.py \
  --csv_input=/content/test_zip/test_labels.csv \
  --output_path=/content/test_zip/test.record

WARNING:tensorflow:From generate_tfrecord.py:107: The name tf.app.run is deprecated.
WARNING:tensorflow:From generate_tfrecord.py:88: The name tf.python_io.TFRecordWriter is deprecated.
W1225 20:27:02.107374 140322520577920 module_wrapper.py:139] From generate_tfrecord.py:107: The name tf.app.run is deprecated.
WARNING:tensorflow:From generate_tfrecord.py:47: The name tf.gfile.GFile is deprecated.
W1225 20:27:02.124306 140322520577920 module_wrapper.py:139] From generate_tfrecord.py:47: The name tf.gfile.GFile is deprecated.
Successfully created the TFRecords: /content/test_zip/test.record
```

```
!pwd
!ls -lh /content/test_zip/*.record
!ls -lh /content/train_zip/*.record

/content/models/research/models/research
-rw-r--r-- 1 root root 6.8M Dec 25 20:27 /content/test_zip/test.record
-rw-r--r-- 1 root root 23M Dec 25 20:27 /content/train_zip/train.record
```

Be sure to store these .record files to somewhere safe. Next, we need to generate a .pbtxt file that defines a mapping between our classes and integers. In the generate_tfrecord.py script, we used the following mapping -

```
def class_text_to_int(row_label):
    if row_label == 'orange':
        return 1
    elif row_label == 'banana':
        return 2
    elif row_label == 'apple':
        return 3
    else:
        return None

label_encodings = {
    "orange": 1,
    "banana": 2,
    "apple": 3
}

f = open("/content/label_map.pbtxt", "w")

for (k, v) in label_encodings.items():
```

```

item = ("item {\n"
        "\tid: " + str(v) + "\n"
        "\tname: '" + k + "'\n"
        "}\n")
f.write(item)

f.close()

!cat /content/label_map.pbtxt

```

```

item {
    id: 1
    name: 'orange'
}
item {
    id: 2
    name: 'banana'
}
item {
    id: 3
    name: 'apple'
}

```

Be sure to save this file as well. Next we will proceed toward training a custom detection model with what we have so far. Follow the steps in [this notebook](#).

In this notebook we will be fine-tuning a **MobileDet** model on the [fruits dataset](#). The original model checkpoints were generated in TensorFlow 1, so we need to stick to a TF 1 runtime. The purpose is to demonstrate the workflow here and not achieve state-of-the-art results. So, please expect unexpected performance for a shorter training schedule. Toward the very end, we will also see how to optimize our fine-tuned model using TensorFlow Lite APIs and run inference with it. This part will be executed on a TF 2 runtime.

As a prerequisite, you should be familiar with the contents of [this notebook](#). It deals with the dataset construction part.

Fetch pre-trained MobileDet model checkpoints and configuration

MobileDet comes in different variants (refer [here](#)). We will be using the `ssdlite_mobiledet_cpu` variant.

TFOD API operates with configuration files to train and evaluate models (the TF 2 release supports eager model execution too). For the purpose of this notebook, I created a configuration file

following instructions from [here](#). Note that I purposefully kept the `num_steps` argument to 2000. Here's a *non-exhaustive* list of the arguments I changed -

- `batch_size: 32`
- `label_map_path` and `input_path` inside `train_input_reader` and `tf_record_input_reader` respectively. The `num_examples` argument inside `eval_config` is set to 117.

▼ Download model checkpoint and config

```
!cd /content/models/research
!wget -q http://download.tensorflow.org/models/object\_detection/ssdlite\_mobiledet\_cpu
!wget -q https://gist.github.com/sayakpaul/9efad54dee957cc55b3adacf992a7a4
```

▼ Untar and verify the file structure of the model checkpoints

```
!cd /content/models/research
!tar -xvf ssdlite_mobiledet_cpu_320x320_coco_2020_05_19.tar.gz
!cp /content/label_map.pbtxt /content/models/research
!cp /content/test_zip/test.record /content/models/research
!cp /content/train_zip/train.record /content/models/research

ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/model.ckpt-400000.data-00000-of-00
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/model.ckpt-400000.index
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/model.ckpt-400000.meta
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/tflite_graph.pbtxt
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/tflite_graph.pb
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/pipeline.config
ssdlite_mobiledet_cpu_320x320_coco_2020_05_19/model.tflite
```

▼ Model training

▼ Start training

Note: This script interleaves both training and evaluation. Before starting the training verify the paths carefully.

```
PIPELINE_CONFIG_PATH="/content/models/research/ssdlite_mobiledet_cpu_320x320_fruits_s
MODEL_DIR="/content/models/research/ssdlite_mobiledet_cpu_320x320_coco_2020_05_19"
```

```

!python object_detection/model_main.py \
  --pipeline_config_path={PIPELINE_CONFIG_PATH} \
  --model_dir={MODEL_DIR} \
  --alsologtostderr

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.52 ^
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.58
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.0
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.57
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.59
INFO:tensorflow:Finished evaluation at 2020-12-25-20:26:00
I1225 20:26:00.477211 140098167961472 evaluation.py:275] Finished evaluation a
INFO:tensorflow:Saving dict for global step 2000: DetectionBoxes_Precision/mAP
I1225 20:26:00.477520 140098167961472 estimator.py:2049] Saving dict for globa
INFO:tensorflow:Saving 'checkpoint_path' summary for global step 2000: /conten

I1225 20:26:00.480742 140098167961472 estimator.py:2109] Saving 'checkpoint_pa
INFO:tensorflow:Performing the final export in the end of training.
I1225 20:26:00.481420 140098167961472 exporter.py:410] Performing the final ex
INFO:tensorflow:Calling model_fn.
I1225 20:26:00.724067 140098167961472 estimator.py:1148] Calling model_fn.
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:03.809365 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:03.881179 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:03.954153 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:04.027084 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:04.103517 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:04.176319 140098167961472 convolutional_box_predictor.py:156] dept
INFO:tensorflow:Done calling model_fn.
I1225 20:26:04.843393 140098167961472 estimator.py:1150] Done calling model_fn
WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/sa
Instructions for updating:
This function will only be available through the v1 compatibility library as t
W1225 20:26:04.843645 140098167961472 deprecation.py:323] From /tensorflow-1.1
Instructions for updating:
This function will only be available through the v1 compatibility library as t
INFO:tensorflow:Signatures INCLUDED in export for Classify: None
I1225 20:26:04.844203 140098167961472 export_utils.py:170] Signatures INCLUDED
INFO:tensorflow:Signatures INCLUDED in export for Regress: None
I1225 20:26:04.844320 140098167961472 export_utils.py:170] Signatures INCLUDED
INFO:tensorflow:Signatures INCLUDED in export for Predict: ['tensorflow/servin
I1225 20:26:04.844394 140098167961472 export_utils.py:170] Signatures INCLUDED
INFO:tensorflow:Signatures INCLUDED in export for Train: None
I1225 20:26:04.844467 140098167961472 export_utils.py:170] Signatures INCLUDED
INFO:tensorflow:Signatures INCLUDED in export for Eval: None
I1225 20:26:04.844529 140098167961472 export_utils.py:170] Signatures INCLUDED
2020-12-25 20:26:04.845001: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:04.845482: I tensorflow/core/common_runtime/gpu/gpu_device.cc
name: Tesla T4 major: 7 minor: 5 memoryClockRate(GHz): 1.59
pciBusID: 0000:00:04.0
2020-12-25 20:26:04.845586: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845621: I tensorflow/stream_executor/platform/default/dso_

```

```

2020-12-25 20:26:04.845647: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845671: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845692: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845711: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845732: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:04.845827: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:04.846229: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:04.846580: I tensorflow/core/common_runtime/gpu/gpu_device.cc

```

The above code block would take approximately **30 minutes** to run (although it depends on the GPU you got if you are running on Colab). If you increase the number of steps it would be even more. After the training was completed I got the following output -

```
I0915 04:48:33.129830 139851326252928 estimator.py:371] Loss for final step: 1.0553685.
```

▼ Export TFLite compatible graph

To export the fine-tuned checkpoints to a TFLite model we first need to export a model graph that is compatible with TFLite. More instructions about this are available [here](#). First, we need to determine which checkpoints to be used to export the graph. Let's first take a look at our MODEL_DIR to get an idea.

```
!ls -lh $MODEL_DIR
```

```

total 300M
-rw-r--r-- 1 root    root    221 Dec 25 20:25 checkpoint
drwxr-xr-x 2 root    root    4.0K Dec 25 20:06 eval_0
-rw-r--r-- 1 root    root    27M Dec 25 20:25 events.out.tfevents.1608926126.0320a
drwxr-xr-x 3 root    root    4.0K Dec 25 20:26 export
-rw-r--r-- 1 root    root    17M Dec 25 19:55 graph.pbtxt
-rw-r--r-- 1 root    root    27M Dec 25 19:55 model.ckpt-0.data-00000-of-00001
-rw-r--r-- 1 root    root    36K Dec 25 19:55 model.ckpt-0.index
-rw-r--r-- 1 root    root    7.4M Dec 25 19:55 model.ckpt-0.meta
-rw-r--r-- 1 root    root    27M Dec 25 20:15 model.ckpt-1321.data-00000-of-00001
-rw-r--r-- 1 root    root    36K Dec 25 20:15 model.ckpt-1321.index
-rw-r--r-- 1 root    root    7.4M Dec 25 20:15 model.ckpt-1321.meta
-rw-r--r-- 1 root    root    27M Dec 25 20:25 model.ckpt-2000.data-00000-of-00001
-rw-r--r-- 1 root    root    36K Dec 25 20:25 model.ckpt-2000.index
-rw-r--r-- 1 root    root    7.4M Dec 25 20:25 model.ckpt-2000.meta
-rw-r----- 1 475825 89939  32M May 19  2020 model.ckpt-400000.data-00000-of-00000
-rw-r----- 1 475825 89939  14K May 19  2020 model.ckpt-400000.index
-rw-r----- 1 475825 89939  9.9M May 19  2020 model.ckpt-400000.meta
-rw-r--r-- 1 root    root    27M Dec 25 20:05 model.ckpt-653.data-00000-of-00001
-rw-r--r-- 1 root    root    36K Dec 25 20:05 model.ckpt-653.index
-rw-r--r-- 1 root    root    7.4M Dec 25 20:05 model.ckpt-653.meta
-rw-r----- 1 475825 89939  16M May 19  2020 model.tflite
-rw-r----- 1 475825 89939  4.6K May 19  2020 pipeline.config

```

```
-rw-r----- 1 475825 89939 17M May 19 2020 tflite_graph.pb
-rw-r----- 1 475825 89939 47M May 19 2020 tflite_graph.pbtxt
```

The checkpoint files with the prefix `model.ckpt-2000` are the ones we would be going with.

```
#Export TFLite compatible graph
#Always verify the paths before running this command.
```

```
!python object_detection/export_tflite_ssd_graph.py \
  --pipeline_config_path=$PIPELINE_CONFIG_PATH \
  --trained_checkpoint_prefix=$MODEL_DIR/model.ckpt-2000 \
  --output_directory=$MODEL_DIR \
  --add_postprocessing_op=true
```



```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tf_slim/layers/
Instructions for updating:
Please use `layer.__call__` method instead.
W1225 20:26:11.416062 140442421712768 deprecation.py:323] From /usr/local/lib/
Instructions for updating:
Please use `layer.__call__` method instead.
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.184998 140442421712768 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.260566 140442421712768 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.334921 140442421712768 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.409476 140442421712768 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.488247 140442421712768 convolutional_box_predictor.py:156] dept
INFO:tensorflow:depth of additional conv before box predictor: 0
I1225 20:26:14.562946 140442421712768 convolutional_box_predictor.py:156] dept
2020-12-25 20:26:14.664051: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.699650: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.700199: I tensorflow/core/common_runtime/gpu/gpu_device.cc
name: Tesla T4 major: 7 minor: 5 memoryClockRate(GHz): 1.59
pciBusID: 0000:00:04.0
2020-12-25 20:26:14.700539: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.702338: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.703908: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.704223: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.714564: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.724804: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.735637: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.735765: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.736402: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.736906: I tensorflow/core/common_runtime/gpu/gpu_device.cc
2020-12-25 20:26:14.748542: I tensorflow/core/platform/profile_utils/cpu_utils
2020-12-25 20:26:14.748726: I tensorflow/compiler/xla/service/service.cc:168]
2020-12-25 20:26:14.748753: I tensorflow/compiler/xla/service/service.cc:176]
2020-12-25 20:26:14.860130: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.860796: I tensorflow/compiler/xla/service/service.cc:168]
2020-12-25 20:26:14.860831: I tensorflow/compiler/xla/service/service.cc:176]
```

```

2020-12-25 20:26:14.861010: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.861537: I tensorflow/core/common_runtime/gpu/gpu_device.cc
name: Tesla T4 major: 7 minor: 5 memoryClockRate(GHz): 1.59
pciBusID: 0000:00:04.0
2020-12-25 20:26:14.861622: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861654: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861680: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861709: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861736: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861760: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861783: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.861864: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.862433: I tensorflow/stream_executor/cuda/cuda_gpu_executo
2020-12-25 20:26:14.862904: I tensorflow/core/common_runtime/gpu/gpu_device.cc
2020-12-25 20:26:14.862971: I tensorflow/stream_executor/platform/default/dso_
2020-12-25 20:26:14.864113: I tensorflow/core/common_runtime/gpu/gpu_device.cc
2020-12-25 20:26:14.864150: I tensorflow/core/common_runtime/gpu/gpu_device.cc
2020-12-25 20:26:14.864165: I tensorflow/core/common_runtime/gpu/gpu_device.cc
2020-12-25 20:26:14.864317: I tensorflow/stream_executor/cuda/cuda_gpu_executo

```

▼ Verify the TFLite compatible graph size

It should have the .pb extension. Be sure to note down the path you would get as the output of code block.

```
!ls -lh $MODEL_DIR/*.pb
```

```
-rw-r----- 1 475825 89939 14M Dec 25 20:26 /content/models/research/ssdlite_mobi
```

Now that we have the graph we can convert it to TensorFlow Lite. Let's shift the runtime to TF 2. To do so, simply restart the Colab runtime.

Optionally see the model losses in TensorBoard (within Colab Notebook)

Note If you trained for 2000 steps only you are likely to see poor numbers in TensorBoard. But as I had mentioned training a SoTA model is not the purpose of this notebook.

```

%tensorflow_version 2.x
%load_ext tensorboard
%tensorboard --logdir /content/models/research/ssdlite_mobiledet_cpu_320x320_coco_202

```


TensorFlow is already loaded. Please restart the runtime to change versions.

TensorBoard

SCALARS

IMAGES

INACTIVE

☐ Show data download links

☐ Ignore outliers in chart scaling

Tooltip sorting
method: default

Smoothing



0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs



TOGGLE ALL RUNS

/content/models/research/ssdlite_
mobiledet_cpu_320x320_coco_2020_
05_19


Filter tags (regular expressions supported)

DetectionBoxes_Precision

6

mAP


tag: DetectionBoxes_Precision/mAP



mAP (large)

tag: DetectionB

Name	Smoothed	Value	Step	Tim
eval_0	0.3678	0.4275	2k	Fri D



mAP (medium)

Export to TFLite

```
#Imports
import tensorflow as tf
print(tf.__version__)
```

```
import os
```

▼ Quantize and serialize

For the purpose of this notebook, we will only be quantizing using the [dynamic-range quantization](#). But you can follow [this notebook](#) if you are interested to try out the other ones like integer quantization and float16 quantization.

As the .pb file we generated in the earlier step is a frozen graph, we need to use `tf.compat.v1.lite.TFLiteConverter.from_frozen_graph` to convert it to TFLite.

The MobileDet checkpoints we used accept 320x320 images, hence the `input_shapes` argument is specified that way. I specified the other arguments following instructions from [here](#). **bold text**

```
model_to_be_quantized = "/content/models/research/ssdlite_mobiledet_cpu_320x320_coco_"
converter = tf.compat.v1.lite.TFLiteConverter.from_frozen_graph(
    graph_def_file=model_to_be_quantized,
    input_arrays=['normalized_input_image_tensor'],
    output_arrays=['TFLite_Detection_PostProcess', 'TFLite_Detection_PostProcess:1', 'T
    input_shapes={'normalized_input_image_tensor': [1, 320, 320, 3]}
)
converter.allow_custom_ops = True
converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_model = converter.convert()

tflite_filename = "fruits_detector" + "_dr" + ".tflite"
open(tflite_filename, 'wb').write(tflite_model)
print(f"TFLite model generated: {tflite_filename}")
!ls -lh $tflite_filename
```

```
TFLite model generated: fruits_detector_dr.tflite
-rw-r--r-- 1 root root 3.6M Dec 25 20:26 fruits_detector_dr.tflite
```

▼ Run inference

```
#Imports
import matplotlib
import matplotlib.pyplot as plt

import cv2
import re
import time
import numpy as np

from PIL import Image
```

▼ TFLite Interpreter and detection utils

Sourced from [here](#).

```
def set_input_tensor(interpreter, image):
    """Sets the input tensor."""
    tensor_index = interpreter.get_input_details()[0]['index']
    input_tensor = interpreter.tensor(tensor_index())[0]
    input_tensor[:, :] = image

def get_output_tensor(interpreter, index):
    """Returns the output tensor at the given index."""
    output_details = interpreter.get_output_details()[index]
    tensor = np.squeeze(interpreter.get_tensor(output_details['index']))
    return tensor

def detect_objects(interpreter, image, threshold):
    """Returns a list of detection results, each a dictionary of object info."""
    set_input_tensor(interpreter, image)
    interpreter.invoke()

    # Get all output details
    boxes = get_output_tensor(interpreter, 0)
    classes = get_output_tensor(interpreter, 1)
    scores = get_output_tensor(interpreter, 2)
    count = int(get_output_tensor(interpreter, 3))

    results = []
    for i in range(count):
        if scores[i] >= threshold:
            result = {
                'bounding_box': boxes[i],
                'class_id': classes[i],
                'score': scores[i]
            }
            results.append(result)
    return results

# Supply a path to download a relevant image
IMAGE_PATH = "https://i.ibb.co/2tsXmCV/image.png"

!wget -q -O image.png $IMAGE_PATH
Image.open('image.png')
```



```
# Load the TFLite model
interpreter = tf.lite.Interpreter(model_path="/content/fruits_detector_dr.tflite")
interpreter.allocate_tensors()
_, HEIGHT, WIDTH, _ = interpreter.get_input_details()[0]['shape']
print(f"Height and width accepted by the model: {HEIGHT, WIDTH}")
```

ValueError

Traceback (most recent call last)

in <ipython input 44: f18b27a61a66> in <module>():

```
# Image preprocessing utils
def preprocess_image(image_path):
    img = tf.io.read_file(image_path)
    img = tf.io.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)
    original_image = img
    resized_img = tf.image.resize(img, (HEIGHT, WIDTH))
    resized_img = resized_img[tf.newaxis, :]
    return resized_img, original_image

# Define the label dictionary and color map
LABEL_DICT = {
    "orange": 1,
    "banana": 2,
    "apple": 3
}

COLORS = np.random.randint(0, 255, size=(len(LABEL_DICT), 3),
                             dtype="uint8")

# Inference utils
def display_results(image_path, threshold=0.3):
    # Load the input image and preprocess it
    preprocessed_image, original_image = preprocess_image(image_path)

    # =====Perform inference=====
    start_time = time.monotonic()
    results = detect_objects(interpreter, preprocessed_image, threshold=threshold)
    print(f"Elapsed time: {(time.monotonic() - start_time)*1000} milliseconds")

    # =====Display the results=====
    original_numpy = original_image.numpy()
    for obj in results:
        # Convert the bounding box figures from relative coordinates
        # to absolute coordinates based on the original resolution
        ymin, xmin, ymax, xmax = obj['bounding_box']
        xmin = int(xmin * original_numpy.shape[1])
        xmax = int(xmax * original_numpy.shape[1])
        ymin = int(ymin * original_numpy.shape[0])
        ymax = int(ymax * original_numpy.shape[0])

        # Grab the class index for the current iteration
        idx = int(obj['class_id'])
        # Skip the background
        if idx >= len(LABEL_DICT):
            continue
```

```

# Draw the bounding box and label on the image
color = [int(c) for c in COLORS[idx]]
cv2.rectangle(original_numpy, (xmin, ymin), (xmax, ymax),
               color, 2)
y = ymin - 15 if ymin - 15 > 15 else ymin + 15
label = "{:}: {:.2f}%".format(LABEL_DICT[idx],
                               obj['score'] * 100)
cv2.putText(original_numpy, label, (xmin, y),
             cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2)

```

```

# return the final image
original_int = (original_numpy * 255).astype(np.uint8)
return original_int

```

```

# Run inference and measure the inference time
resultant_image = display_results("/content/image.png", threshold=0.3)
Image.fromarray(resultant_image)

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-73-a1d5f786e226> in <module>()
      1 # Run inference and measure the inference time
----> 2 resultant_image = display_results("/content/image.png", threshold=0.3)
      3 Image.fromarray(resultant_image)

```

```

----- 1 frames -----
<ipython-input-70-e89900eb19d2> in preprocess_image(image_path)
      5 img = tf.image.convert_image_dtype(img, tf.float32)
      6 original_image = img
----> 7 resized_img = tf.image.resize(img, (HEIGHT, WIDTH))
      8 resized_img = resized_img[tf.newaxis, :]
      9 return resized_img, original_image

```

NameError: name 'HEIGHT' is not defined

SEARCH STACK OVERFLOW

****Note**** that you might see some unexpected results because the annotations in the tr

```

# Define the label dictionary and color map
LABEL_DICT = {
    "orange": 1,
    "banana": 2,
    "apple": 3
}

```

```

COLORS = np.random.randint(0, 255, size=(len(LABEL_DICT), 3),
                             dtype="uint8")

```

```

# Inference utils

```

```

def display_results(image_path, threshold=0.3):
    # Load the input image and preprocess it
    preprocessed_image, original_image = preprocess_image(image_path)

    # =====Perform inference=====
    start_time = time.monotonic()
    results = detect_objects(interpreter, preprocessed_image, threshold=threshold)
    print(f"Elapsed time: {(time.monotonic() - start_time)*1000} milliseconds")

    # =====Display the results=====
    original_numpy = original_image.numpy()
    for obj in results:
        # Convert the bounding box figures from relative coordinates
        # to absolute coordinates based on the original resolution
        ymin, xmin, ymax, xmax = obj['bounding_box']
        xmin = int(xmin * original_numpy.shape[1])
        xmax = int(xmax * original_numpy.shape[1])
        ymin = int(ymin * original_numpy.shape[0])
        ymax = int(ymax * original_numpy.shape[0])

        # Grab the class index for the current iteration
        idx = int(obj['class_id'])
        # Skip the background
        if idx >= len(LABEL_DICT):
            continue

        # Draw the bounding box and label on the image
        color = [int(c) for c in COLORS[idx]]
        cv2.rectangle(original_numpy, (xmin, ymin), (xmax, ymax),
                      color, 2)
        y = ymin - 15 if ymin - 15 > 15 else ymin + 15
        label = "{:}: {:.2f}%".format(LABEL_DICT[idx],
                                      obj['score'] * 100)
        cv2.putText(original_numpy, label, (xmin, y),
                   cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2)

    # return the final image
    original_int = (original_numpy * 255).astype(np.uint8)
    return original_int

# Run inference and measure the inference time
resultant_image = display_results("/content/image.png", threshold=0.3)
Image.fromarray(resultant_image)

```

Note that you might see some unexpected results because the annotations in the training dataset are faulty at places. Due to this the model training can suffer a lot.

